

XGBoost: an extremely boosting method



Prepared by:

A. Kazarian, D. Senkevich, J. Grabenhofer
N. Karetnikov, R. Köfer, I. Kravchenko

07.04.2022



- 1. Introduction
- 2. XGBoost Regression
- 3. XGBoost Classification: Intuition
- 4. Implementation of XGBoost in R
- 5. Comparison of performance
- 6. Distinctive features of XGBoost

1. Introduction

▲ Some **selling points** of XGBoost before we start:


- XGboost is like generalized boosting - but EXTREME!!
- XGboost is widely used in the winning solutions of Kaggle and KGG Cup

Original paper: Chen, T., Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. arXiv:1603.02754.

2. XGBoost Regression: Intuition

▲ Let's follow the algorithm:

1. Calculate the residuals
2. Make the first split of residuals and calculate $Similarity\ score = \frac{(\sum Residuals)^2}{Number\ of\ Residuals + \lambda}$
3. Find $Gain = Left_{Similarity} + Right_{Similarity} - Root_{Similarity}$. Leave the first split in a form which maximizes gain
4. Continue splitting. Default depth is 6 levels
5. Prune the trees: if $Gain - \gamma$ (3 by default) < 0, we prune
6. Calculate $Output = \frac{\sum Residuals}{Number\ of\ Residuals + \lambda}$
7. Find $Predicted\ values = Default\ (0.5\ by\ default) + Learning\ rate \times Output$
8. Go to step 1: calculate new residuals from the new predicted values



2. XGBoost Regression: Maths behind

Loss function $L(y, p) = \frac{1}{2}(y - p)^2$

Expression to be minimized $[\sum_{i=1}^n L(y_i, p_i)] + \gamma T + \frac{1}{2}\lambda O^2$

In other words, $\min[\sum_{i=1}^n L(y_i, p^0 + O)] + \frac{1}{2}\lambda O^2]$

Using Taylor second order approx, we find that $O = \frac{-(g_1+g_2+\dots+g_n)}{(h_1+h_2+\dots+h_n+\lambda)}$,

where g is the first order deriv.: $-(y-p)$ – residuals

and h is the second order deriv.: $1 \times n$ – number of residuals

Now look at the formulas from the previous 'intuition' slide

3. XGBoost Classification: Intuition (1)

▲ Let's follow the algorithm:


1. Calculate the residuals based on initial guess of 0.5
2. Make the first split (average of last two observations) and calculate

$$Similarityscore = \frac{(\sum Residuals)^2}{\sum [Previous\ probability \times (1 - Previous\ probability)] + \lambda}$$
 of each leaf
3. Calculate $Gain = Left_{Similarity} + Right_{Similarity} - Root_{Similarity}$
4. Continue splitting with highest gain split. Default depth is 6 levels
5. Prune the trees by calculating $Gain - \gamma$ (3 by default). If negative we prune the branch
6. Calculate $Output = \frac{\sum Residuals}{\sum [Previous\ probability \times (1 - Previous\ probability)] + \lambda}$
7. Find $Predicted\ values = \log\ odds\ of\ default + Learning\ rate\ (0.3\ by\ default) \times Output$
8. Go to step 1: calculate new residuals from the new predicted values


Different Loss function $L(y, p) = -[y \log(p) + (1 - y) \log(1 - p)]$

3. XGBoost Classification: Intuition (2)


1. Calculate residuals




2. Split



3. Calculate similarities and gain



4. Repeat for other splits




5. Pruning


Gain - γ

{ If positive, then do not prune.
If negative, then prune.

6. Calculating output



7. Predicting values



$\log(\text{odds})$ Prediction = 0 + (0.3 x -2) = -0.6

Probability = $\frac{e^{\log(\text{odds})}}{1 + e^{\log(\text{odds})}}$

4. Implementation of XGBoost in R

Training of the model

NOTE: XGBoost works only with numeric values

```
data = load(file = "german.rda")
dummy = dummyVars("~.", data = german[, -c(2, 5, 13, 21)])
newdata = data.frame(predict(dummy, german[, -c(2, 5, 13, 21)]))
data = cbind(newdata, german[, c(2, 5, 13, 21)])
y = recode(data$Class, 'good'=1, 'bad'=0)

xgb = xgboost(data = data.matrix(data[, -c(length(data))]),
              label = y, eta = 0.3,
              max_depth = 15,
              nround=25, objective = "binary:logistic" ,
              lambda = 1,
              )
```


Making predictions

```
predict(xgb, data.matrix(data[4,-c(length(data))]))  
## [1] 0.8940502
```

5. Comparison of performance

Three algorithms:

- XGBoost
- LightGBM
- CatBoost

Defining functions(1/3) – Xgboost

```
give_xgboost_AUC <- function(X_train,X_test,y_train, y_test){
  xgb = xgboost(data = data.matrix(X_train[,-c(length(X_train))]),
                objective = "binary:logistic" ,
                label = y_train,nrounds=25)
  vec.prob=predict(xgb, data.matrix(X_test[,-c(length(X_test))]))

  pred <- prediction(vec.prob, y_test)
  auc = performance(pred, measure = "auc")
  return(as.numeric(auc@y.values) )
}
```

Defining functions(2/3) – Lightgbm

```

give_lightgbm_AUC <- function(X_train,X_test,y_train, y_test){

fit <- lightgbm(data = data.matrix(X_train[,-c(length(X_train))]),
               label = y_train,
               objective = "binary",nrounds=25)

vec.prob=predict(fit, data.matrix(X_test[,-c(length(X_test))]))
pred <- prediction(vec.prob, y_test)
auc = performance(pred, measure = "auc")
return(as.numeric(auc@y.values)  )
}

```

Defining functions (3/3) – CatBoost

```
give_catboost_AUC <- function(X_train,X_test,y_train, y_test){
  train_pool <- catboost.load_pool(
    data = X_train[,-c(length(X_train))], label = y_train)
  test_pool <- catboost.load_pool(
    data = X_test[,-c(length(X_test))], label = y_test)

  params <- list(iterations=25, loss_function='CrossEntropy',
    eval_metric='CrossEntropy',
    use_best_model=TRUE)
  fit <- catboost.train(train_pool,test_pool, params)

  vec.prob= catboost.predict(fit, test_pool)
  pred <- prediction(vec.prob, y_test)
  auc = performance(pred, measure = "auc")
  return(as.numeric(auc@y.values))
}
```

Loop for comparing the performance

```
xgboost_AUC<-lightgbm_AUC<-catboost_AUC<-NULL
for (i in (1:1000)){
  vec <- sample(dim(data)[1],700)

  X_train <- data[vec,]
  X_test <- data[-vec,]
  y_train <- y[vec]
  y_test <- y[-vec]

  xgboost_AUC=c(xgboost_AUC,
                give_xgboost_AUC(X_train,X_test,y_train, y_test))
  lightgbm_AUC=c(lightgbm_AUC,
                  give_lightgbm_AUC(X_train,X_test,y_train, y_test))
  catboost_AUC=c(catboost_AUC,
                  give_catboost_AUC(X_train,X_test,y_train, y_test))}
```

Results of the comparison

```
print(c(mean (xgboost_AUC),mean (lightgbm_AUC),mean (catboost_AUC) ))  
## [1] 0.7726709 0.7728125 0.7663909  
  
print(sum(xgboost_AUC>lightgbm_AUC))  
## [1] 500  
  
print(sum(xgboost_AUC>catboost_AUC))  
## [1] 600
```

6. Distinctive features of XGBoost

When the dataset gets large, XGBoost becomes handy

- Approximate Greedy Algorithm
- Parallel Learning
- Weighted Quantile Sketch
- Sparsity-Aware Split Finding
- Cache-Aware Access
- Blocks for Out-of-Core Computation

Grand finale

Thank you for attention!

(Invite us for a coffee afterwards if you want to speak about shap values)